

資料庫管理 HW04

B12508026 戴偉璿

November 16, 2025

1. To check if PostgreSQL can avoid dirty read, I design two transactions:

- Transaction A: Update balance to 999 of account_id 1

```
1      begin;
2      update accounts set balance = 999 where account_id = 1;
3      commit;
```

- Transaction B: Read the record.

```
1      begin; select * from accounts where account_id = 1; commit;
2
3
```

The execution steps are as follows:

- (a) Transaction A begins.
- (b) Transaction A updates balance to 999 of account_id 1, but does not commit yet.
- (c) Transaction B begins.
- (d) Transaction B reads the record of account_id 1.
- (e) Transaction B gets the old balance (not 999), which means dirty read is avoided.
- (f) Transaction B commits.
- (g) Transaction A commits.
- (h) Transaction B begins.
- (i) Transaction B reads the record of account_id 1.
- (j) Transaction B gets the new balance (999) after Transaction A commits.
- (k) Transaction B commits.

Following are the screenshots of each step. Left panel shows Transaction A, right panel shows Transaction B. Figure 1 shows the original status of the accounts table, we can see the balance of account_id 1 is 1000. Figure 2 shows Transaction A updates balance to

999 of account_id 1 but does not commit yet, so that Transaction B still reads the old balance (1000). Figure 3 shows Transaction A commits, and then Transaction B reads the new balance (999) of account_id 1. You can determine the execution order by the system time shown in the top of each figure.

```
11月13日 週四 22:00:35 zh
josh@josh-ubuntu:~ postgres=# select * from accounts;
account_id | balance
-----+-----
2 |      10
1 |    1000
(2 rows)
postgres=#[[REDACTED]
```

Figure 1: Orginal status of the accounts table

```
11月13日 週四 22:02:23 zh
josh@josh-ubuntu:~ postgres=# select * from accounts;
account_id | balance
-----+-----
2 |      10
1 |    1000
(2 rows)

postgres=# begin; select * from accounts where account_id = 1; commit;
BEGIN
account_id | balance
-----+-----
1 |    1000
(1 row)

COMMIT
postgres=#[[REDACTED]
```

Figure 2: Transaction A updates balance to 999 of account_id 1, but does not commit yet

```

11月13日 週四 22:02:45 UTC
josh@josh-ubuntu:~ postgres=# begin; select * from accounts where account_id = 1; commit;
BEGIN
account_id | balance
-----+-----
1          |    1000
(1 row)

COMMIT
postgres=# begin; update accounts set balance = 999 where account_id = 1;
BEGIN
UPDATE 1
postgres=# commit;
COMMIT
postgres=#

```

```

11月13日 週四 22:02:45 UTC
josh@josh-ubuntu:~ postgres=# begin; select * from accounts where account_id = 1; commit;
BEGIN
account_id | balance
-----+-----
1          |    1000
(1 row)

COMMIT
postgres=# begin; select * from accounts where account_id = 1; commit;
BEGIN
account_id | balance
-----+-----
1          |     999
(1 row)

COMMIT
postgres=#

```

Figure 3: Transaction A commits, Transaction B reads the new balance (999) of account_id 1

With the experiments above, we can see that PostgreSQL can avoid dirty read.

2. (a) A conflict occurs when two transactions access the same data item and at least one of the accesses is a write operation. For item X , three conflicts occurs between $\{O_{11}, O_{23}\}$, $\{O_{17}, O_{21}\}$, $\{O_{17}, O_{23}\}$; for item Y , there are no conflicts; for item Z , two conflicts occurs between $\{O_{15}, O_{24}\}$, $\{O_{15}, O_{26}\}$.
- (b) The serial schedule of $T_2 \rightarrow T_1$ is as follows:

$$O_{21} \rightarrow O_{23} \rightarrow O_{24} \rightarrow O_{26} \rightarrow O_{11} \rightarrow O_{12} \rightarrow O_{13} \rightarrow O_{15} \rightarrow O_{17}.$$

To analyze the conflicting operation pairs, we can discuss by data items:

For item X , the operation order must obey: $O_{21} \prec O_{23} \prec O_{17}$, $O_{23} \prec O_{11} \prec O_{17}$

For item Y , there are no conflicts, so there is no constraint.

For item Z , the operation order must obey: $O_{24} \prec O_{26} \prec O_{15}$

Therefore, one such conflict-equivalent non-serial schedule is:

$$O_{21} \rightarrow O_{23} \rightarrow O_{11} \rightarrow O_{24} \rightarrow O_{26} \rightarrow O_{12} \rightarrow O_{13} \rightarrow O_{15} \rightarrow O_{17}.$$

- (c) The serial schedule of $T_1 \rightarrow T_2$ is as follows:

$$O_{11} \rightarrow O_{12} \rightarrow O_{13} \rightarrow O_{15} \rightarrow O_{17} \rightarrow O_{21} \rightarrow O_{23} \rightarrow O_{24} \rightarrow O_{26}.$$

Consider the last operation of T_1 and the first operation of T_2 , we have $O_{17} \prec O_{21}$ because O_{17} is the write operation and they access the same data item X .

To maintain the order within the transactions, O_{17} must be the last operation in the schedule and O_{21} must be the first operation in the schedule. From the analyze above, O_{21} must after O_{17} . Therefore, there is no such conflict-equivalent non-serial schedule.