

Fundamentals of Biomedical Image Processing HW 1

B12508026 戴偉璿

September 30, 2025

1 Theorem questions

1. (a) The picture size is 1024×1024 pixels, and each pixel contains 256 intensity levels(8 bits,i.e., 1 byte). Therefore, the total data size is $1024 \times 1024 \times 1 = 1,048,576$ bytes. Since each packet requires 10 bits to transmit 1 byte of data, the total number of bits to be transmitted is 10,485,760 bits. With the baud rate of 56k, the transmission time is $10,485,760 \div 56000 \approx 187.25(sec) \approx 3.12$ minutes.
- (b) With the baud rate of 3000k, the transmission time is $10,485,760 \div 3,000,000 \approx 3.495(sec) \approx 0.058$ minutes.
2. $V = \{1, 2\}$, so the graph would be Fig. 1:

				q
	3	2	3	2
	2	2	1	3
	1	3	3	2
p	2	2	3	4

Figure 1: The grid graph with $V = \{1, 2\}$ (gray cells are passable)

- (a) shortest-4: Only up, down, left, right movements are allowed. So the path would be:

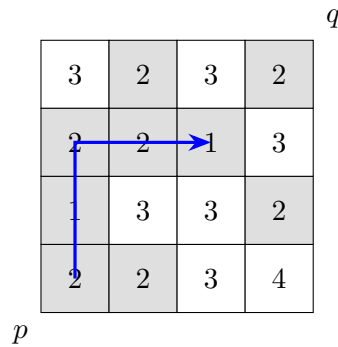


Figure 2: No 4-connected path (q is 4-blocked)

As the Fig. 2 shows, there's no path from last step (3,3) to node $q(4,4)$, so the length is **infinity**.

- (b) shortest-8: Up, down, left, right, and diagonal movements are allowed. So the path would be:

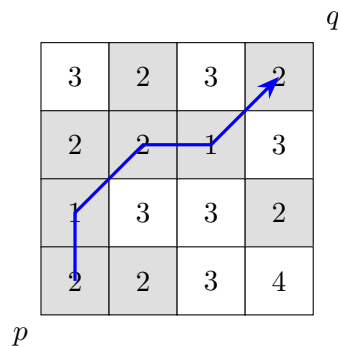


Figure 3: shortest-8 path (length = 4)

As the Fig. 3 shows, it takes 4 steps to reach from $p(1,1)$ to $q(4,4)$, so the length is **4**.

(c) shortest-m: The path can only move to the 8-neighbors with the minimum value.
So the path would be:




Figure 4: shortest-m path (length = 5)

As the Fig. 4 shows, it takes 5 steps to reach from $p(1, 1)$ to $q(4, 4)$,
so the length is **5**.

2 Programming exercises

All code can be found at <https://github.com/W-X-Dai/tex/tree/main/imgProc/HW1/> code.

1. To draw a triangle, we can find its three vertices, and then use the line-drawing algorithm to connect them. The result is Fig. 5:




Figure 5: The result of drawing an equilateral triangle

Following is the code:

```
1 import os
2 import cv2
3 import numpy as np
4
5 os.makedirs("result/p1", exist_ok=True)
6
7 img = cv2.imread("src/fig1.bmp", cv2.IMREAD_GRAYSCALE)
8 if img is None:
9     print("Error: Unable to load image.")
10    exit()
11
12 """Calculating the vertices coordinates of an equilateral triangle."""
13 top = (150, 150)
14 h = int((3*0.5 / 2) * 200)
15 left = (150 - 100, 150 + h)
16 right = (150 + 100, 150 + h)
17
18 pts = np.array([top, left, right], np.int32)
19
20 cv2.fillPoly(img, [pts], 255)
21
22 cv2.imwrite("result/p1/output.bmp", img)
23 print("Result image saved to result/p1/output.bmp")
24
25 cv2.imshow("Triangle", img)
26 cv2.waitKey(0)
27 cv2.destroyAllWindows()
```

2. To transfer the gray levels of the original image to assigned levels, we can use the formula:

$$\text{new_pixel} = \left\lfloor \frac{\text{old_pixel} \cdot \text{levels}}{256} \right\rfloor + \frac{128}{\text{levels}}$$

The result is in Fig. 6:




Figure 6: Gray images from 2 to 256 (256 is the original image)

Following is the code:

```

1  import os
2  import cv2
3  import numpy as np
4
5  os.makedirs("result/p2", exist_ok=True)
6
7  def reduce_gray_levels(img, nLevels=2):
8      if nLevels < 2 or nLevels > 256 or (nLevels & (nLevels - 1)):
9          raise ValueError("invalid input: nLevels must be a power of 2 and between 2 and 256.")
10
11     delta = 256 // nLevels
12     quantized = (img // delta) * delta + delta // 2
13     return quantized.astype(np.uint8)
14
15  img = cv2.imread("src/fig1.bmp", cv2.IMREAD_GRAYSCALE)
16  if img is None:
17      raise FileNotFoundError("Cannot open fig1.bmp")
18
19  nLevels = int(input("Enter the number of gray levels (should be a power of 2 and between 2 and 256,
    ↪ e.g., 2,4,8,...,256): "))
20
21  out = reduce_gray_levels(img, nLevels)

```

```

22
23 cv2.imwrite(f"result/p2/gray_{nLevels}.bmp", out)
24 print(f"Result image saved to result/p2/gray_{nLevels}.bmp")
25
26 cv2.imshow(f"{nLevels} levels", out)
27 cv2.waitKey(0)
28 cv2.destroyAllWindows()

```

3. To rotate the image by a specified angle, we can use the rotation matrix to calculate the new coordinates of each pixel. I once implemented a bilinear interpolation algorithm and then modified it into a nearest neighbor interpolation. The result is shown in Fig. 7.




Figure 7: The result of rotation

Following is the code:

```

1 import os
2 import cv2
3 import numpy as np
4 import math
5 import matplotlib.pyplot as plt
6
7 """This is used to be a bilinear interpolation function I've finished in the past, but I changed to
  ↪ nearest neighbor interpolation for simplicity."""
8 def bi_affine(img, A):
9     img_out = np.zeros_like(img)
10     inv = np.linalg.inv(A)

```

```

11
12     h, w = img.shape
13     for i in range(h):          # row
14         for j in range(w):      # col
15             origin_T = np.array([[j], [i], [1]])
16             inv_T = inv @ origin_T
17
18             c, r = inv_T[0].item(), inv_T[1].item()
19
20             # clamp
21             r = min(max(r, 0), h-1)
22             c = min(max(c, 0), w-1)
23
24             # nearest neighbor
25             r_nn, c_nn = int(round(r)), int(round(c))
26             img_out[i, j] = img[r_nn, c_nn]
27
28     return img_out
29
30 def get_rotation_matrix(h, w, angle_deg):
31     cx, cy = w/2, h/2
32     theta = math.radians(angle_deg)
33     cos_t, sin_t = math.cos(theta), math.sin(theta)
34
35     A = np.array([
36         [cos_t, sin_t, (1-cos_t)*cx-sin_t*cy],
37         [-sin_t, cos_t, (1-cos_t)*cy+sin_t*cx],
38         [0, 0, 1]
39     ], dtype=np.float32)
40
41     return A
42
43 if __name__ == '__main__':
44     os.makedirs("result/p3", exist_ok=True)
45
46     img = cv2.imread('src/Fig2.bmp', cv2.IMREAD_GRAYSCALE)
47     if img is None:
48         raise FileNotFoundError("Cannot find Fig2.bmp")
49
50     h, w = img.shape
51
52     angles = [15, 30, 45, 60, 90]
53     results = [img]
54
55     for angle in angles:
56         """minus angle for clockwise rotation"""
57         A = get_rotation_matrix(h, w, -angle)
58         rotated = bi_affine(img, A)
59
60         cv2.imwrite(f"result/p3/rot_{angle}.bmp", rotated)
61         results.append(rotated)
62
63     print("Result images saved to result/p3/")
64
65     plt.figure(figsize=(12, 8))
66     titles = ["Original"] + [f"{a}° CW" for a in angles]
67
68     for idx, (title, im) in enumerate(zip(titles, results), 1):
69         plt.subplot(2, 3, idx)

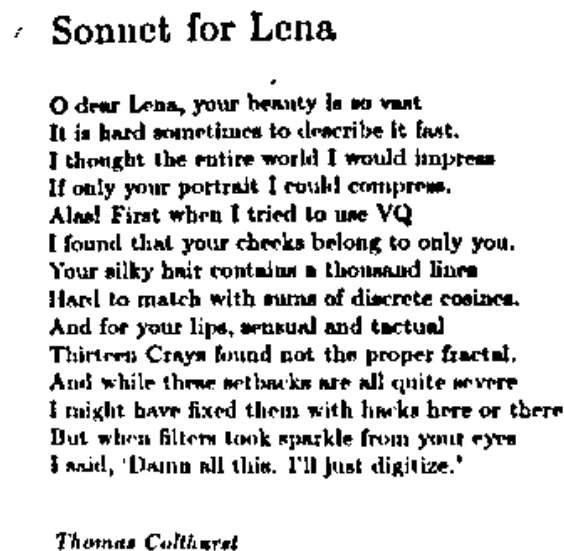
```

```

70     plt.imshow(im, cmap="gray")
71     plt.title(title)
72     plt.axis("off")
73
74     plt.tight_layout()
75     plt.savefig("result/p3/combined_grid.png")
76     plt.show()

```

4. Recovering text from an image containing a strong illumination gradient is a challenging task. I first applied CLAHE to enhance the local contrast of the image and reduce the effect of uneven illumination. Then, I used adaptive thresholding to binarize the image and recover the text. The result is shown in Fig. 8:



Sonnet for Lena

O dear Lena, your beauty is so vast
 It is hard sometimes to describe it fast.
 I thought the entire world I would impress
 If only your portrait I could compress.
 Alas! First when I tried to use VQ
 I found that your cheeks belong to only you.
 Your silky hair contains a thousand lines
 Hard to match with sums of discrete cosines.
 And for your lips, sensual and tactual
 Thirteen Crays found not the proper fractal.
 And while these setbacks are all quite severe
 I might have fixed them with hacks here or there
 But when filters took sparkle from your eyes
 I said, 'Damn all this. I'll just digitize.'

Thomas Colthurst

Figure 8: The result of text recovery

Following is the code:

```

1  import cv2
2  import os
3  import numpy as np
4
5  os.makedirs("result/p4", exist_ok=True)
6
7  img = cv2.imread("src/Fig3.GIF", cv2.IMREAD_GRAYSCALE)
8  if img is None:
9      raise FileNotFoundError("Cannot open Fig3.GIF")
10

```

```

11 clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8,8))
12 img_eq = clahe.apply(img)
13
14 """A good result is obtained with block size = 25 and C = 10."""
15 binary = cv2.adaptiveThreshold(img_eq, 255,
16                               cv2.ADAPTIVE_THRESH_GAUSSIAN_C,
17                               cv2.THRESH_BINARY,
18                               25, 10)
19 cv2.imwrite("result/p4/text_recovered.bmp", binary)
20
21 print("Saved result to result/p4/text_recovered.bmp")

```

As the picture shows, though most of the noise is removed, but the result of the text recovery is still not very good. Some of the text structures are broken. To address this, I tried to train a Denoisy UNet to further improve the image quality.

Denoisy UNet: If you are interested in the code of Denoisy UNet, you can find it at https://github.com/W-X-Dai/text_recovery.

For the training data, I generated 20000 clean samples, each sample contains a white background with black text. Then, I added Gaussian noise and illumination gradient to these clean samples to create noisy samples. Finally, I use my code in this homework to simulate the result of CLAHE and adaptive thresholding, and use these processed images as the input of the Denoisy UNet. Fig. 9 is an example of noisy data:




Figure 9: An example of the noisy data

And Fig. 10 is an example of the train result after 100 epochs: (Left is the broken input, middle is the output of the Denoisy UNet, right is the ground truth)

iAUahIRbCsCufzYmB PRjsKYKciMmldkGRZk tWOVpwzKgnzBQEfKLHrOjJRbKI UaEgflpcYgBgEjgcjDBBAQioDfu yAkpVmuRGoGxnmItWHXXIcRWI CwglIHlPhtKNRvAeFefTpgsABvRV TuWgbFbSroZMSQGOc ZvmaWF rwJSVncjUyrCiRuTdOFaXKCOIHx hbcarZifNWQICrZEYAavSjzbAbG ezCoaZjsxjUPRpJUkrVBNTliP Ab evnmAmelgzRpqZyTGzpyS wNStSwC kumlNtQgnFJhkH eBU	iAUahIRbCsCufzYmB PRjsKYKciMmldkGRZk tWOVpwzKgnzBQEfKLHrOjJRbKI UaEgflpcYgBgEjgcjDBBAQioDfu yAkpVmuRGoGxnmItWHXXIcRWI CwglIHlPhtKNRvAeFefTpgsABvRV TuWgbFbSroZMSQGOc ZvmaWF rwJSVncjUyrCiRuTdOFaXKCOIHx hbcarZifNWQICrZEYAavSjzbAbG ezCoaZjsxjUPRpJUkrVBNTliP Ab evnmAmelgzRpqZyTGzpyS wNStSwC kumlNtQgnFJhkH eBU	iAUahIRbCsCufzYmB PRjsKYKciMmldkGRZk tWOVpwzKgnzBQEfKLHrOjJRbKI UaEgflpcYgBgEjgcjDBBAQioDfu yAkpVmuRGoGxnmItWHXXIcRWI CwglIHlPhtKNRvAeFefTpgsABvRV TuWgbFbSroZMSQGOc ZvmaWF rwJSVncjUyrCiRuTdOFaXKCOIHx hbcarZifNWQICrZEYAavSjzbAbG ezCoaZjsxjUPRpJUkrVBNTliP Ab evnmAmelgzRpqZyTGzpyS wNStSwC kumlNtQgnFJhkH eBU
--	--	--

Figure 10: An example of the training data

Then I applied this model to the result of CLAHE and adaptive thresholding, and got the final result shown in Fig. 11:

Sonnet for Lena

O drer Lena, your beauty is so vast
 It is hard sometimes to describe it last
 I thought the entire world I could impress
 If only your portrait I could compress
 Alas! First what I tried to use VQ
 I found that your cheeks belong to only you
 Your silky hair contains a thousand lines
 Hard to match with sums of discrete cosines
 And for your lips, sensual and tactile
 Thirteen Cmvx found not the proper fractal
 And still their webarks are quite worse
 I might have fixed them with hark-hark or three
 But when fifteen took sparkle from your eyes
 I said, Damn all this, I'll just digitize

Tao-mao Colfaxerzi

Figure 11: The final result after using Denoisy UNet

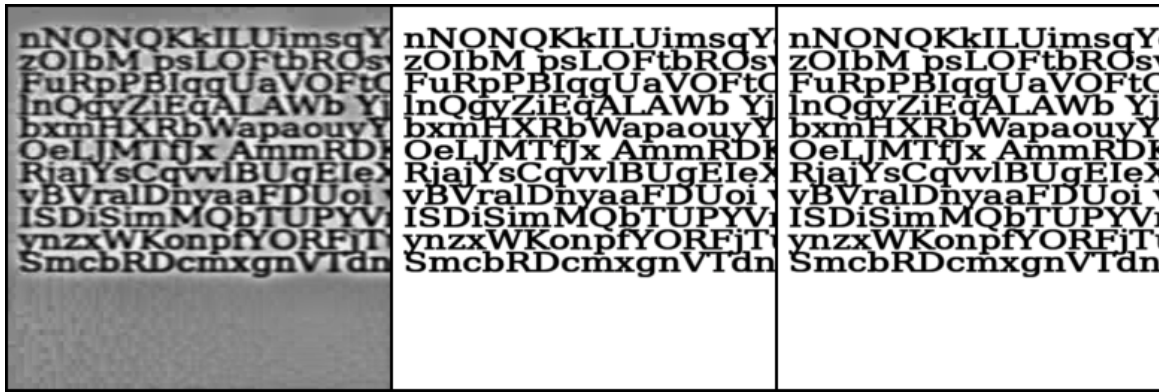


Figure 12: New training data without adaptive thresholding

But as the figure shows, the result is still not very good. Some of the text structures are distorted. The structure broken of adaptive thresholding is still too severe. I suddenly thought that the result of CLAHE is already quite good for this model (because the background color is uniform). Then I created another set of training data and trained the new model. Fig. 12 shows the result after 100 epochs of training. Left is the input (only CLAHE, no adaptive thresholding), middle is the output of the Denoisy UNet, right is the ground truth. And then I applied this new model to the target image(after CLAHE only), and got a much better result shown in Fig. 13:

Sonnet for Lena

O dear Lena, your beauty is so vast
It is hard sometimes to describe it fast.
I thought the entire world I would impress
If only your portrait I could compress.
Alas! First when I tried to use VQ
I found that your cheeks belong to only you.
Your silky hair contains a thousand lines
Hard to match with sums of discrete cosines.
And for your lips, sensual and tactual
Thirteen Crays found not the proper fractal.
And while these setbacks are all quite severe
I might have fixed them with hacks here or there
But when filters took sparkle from your eyes
I said, 'Damn all this. I'll just digitize.'

Thomas Collier

Figure 13: The final result after using the new Denoisy UNet