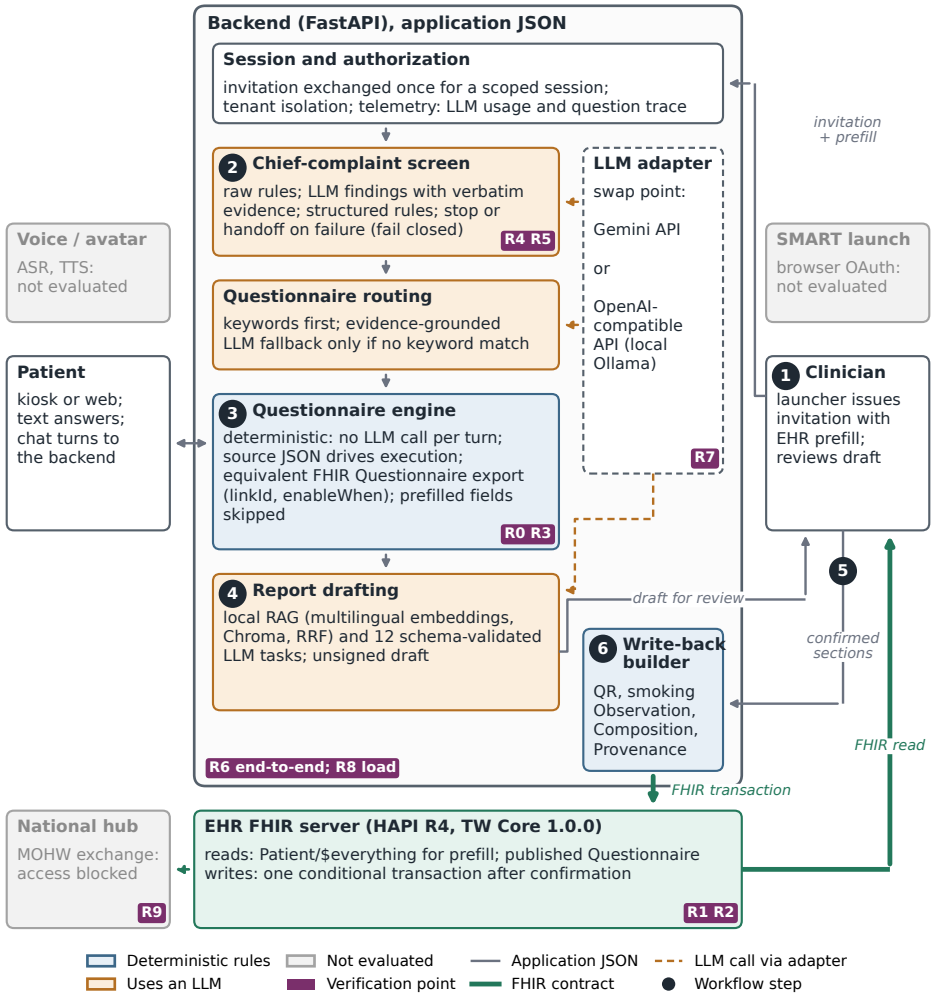
Backend (FastAPI), application JSON
Session and authorization
invitation exchanged once for a scoped session; tenant isolation; telemetry: LLM usage and question trace
invitation + prefill
2 Chief-complaint screen
raw rules; LLM findings with verbatim evidence; structured rules; stop or handoff on failure (fail closed)
R4 R5
LLM adapter
swap point:
Gemini API
or
OpenAI-compatible API (local Ollama)
R7
Questionnaire routing
keywords first; evidence-grounded LLM fallback only if no keyword match
Voice / avatar
ASR, TTS: not evaluated
SMART launch
browser OAuth: not evaluated
Patient
kiosk or web; text answers; chat turns to the backend
3 Questionnaire engine
deterministic: no LLM call per turn; source JSON drives execution; equivalent FHIR Questionnaire export (linkId, enableWhen); prefilled fields skipped
R0 R3
1 Clinician
launcher issues invitation with EHR prefill; reviews draft
5
4 Report drafting
local RAG (multilingual embeddings, Chroma, RRF) and 12 schema-validated LLM tasks; unsigned draft
draft for review
confirmed sections
6 Write-back builder
QR, smoking Observation, Composition, Provenance
FHIR read
R6 end-to-end; R8 load
FHIR transaction
National hub
MOHW exchange: access blocked
R9
EHR FHIR server (HAPI R4, TW Core 1.0.0)
reads: Patient/$everything for prefill; published Questionnaire
writes: one conditional transaction after confirmation
R1 R2
Deterministic rules
Uses an LLM
Not evaluated
Verification point
Application JSON
FHIR contract
LLM call via adapter
Workflow step